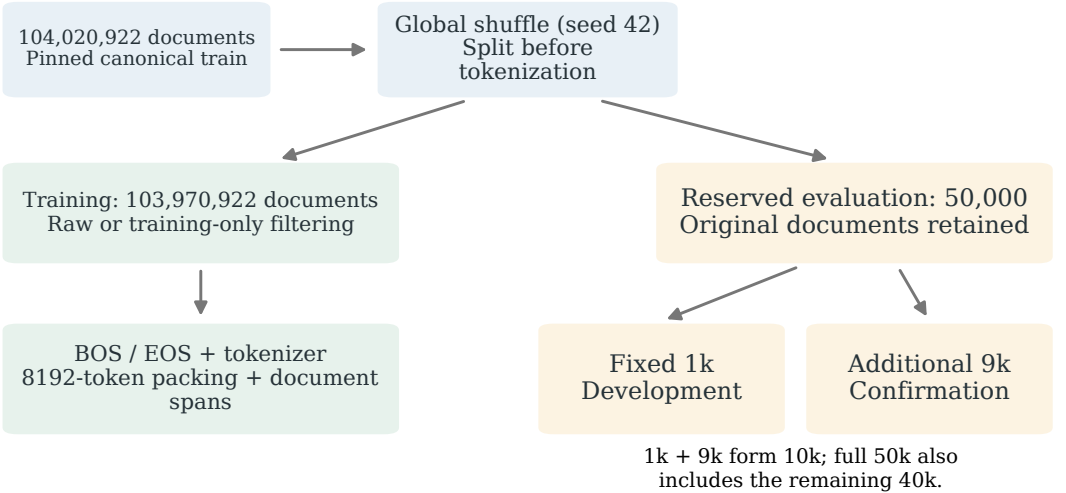
104,020,922 documents
Pinned canonical train
Global shuffle (seed 42)
Split before tokenization
Training: 103,970,922 documents
Raw or training-only filtering
Reserved evaluation: 50,000
Original documents retained
BOS / EOS + tokenizer
8192-token packing + document spans
Fixed 1k
Development
Additional 9k
Confirmation
1k + 9k form 10k; full 50k also includes the remaining 40k.